

CHATBOT FOR INTENT RECOGNITION

Presented by
Adina Dingankar
Rehapriadarsini Manikandasamy
Venkata Gangadhar Naveen Palaka





INTENT RECOGNITION (IR)

- ❑ Intent recognition is sometimes called as **intent classification** which is the task of taking a written or spoken input, and classifying it based on what the user wants to achieve
- ❑ Intent recognition works through the process of providing examples of text alongside their intents to a machine learning (ML) model

APPLICATIONS: Chatbots, customer support, sales prospecting, etc.,



CHATBOTS USE IR?

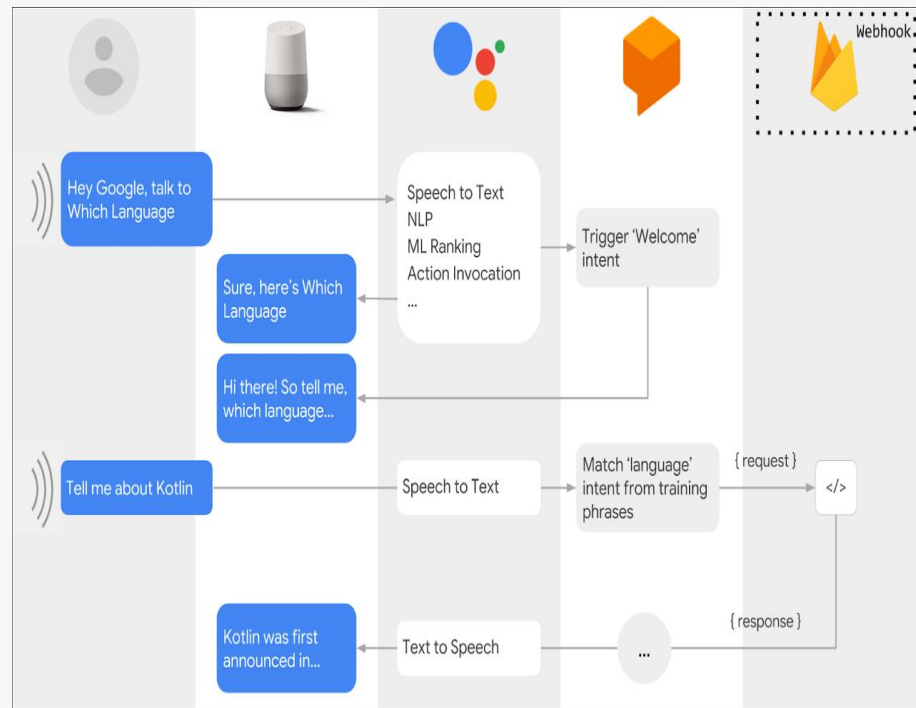
- ❑ Intent recognition is a critical feature in chatbot architecture that determines user's aim in starting any conversation
- ❑ NLP allows the chatbot to understand the user's message, and machine learning classification algorithms to classify the message based on the training data and deliver the correct response
- ❑ The chatbots' intent detection component helps to identify what general task or goal the user is trying to accomplish to handle the conversation with different strategies



GOOGLE & IR

- **Google's Dialogflow** - Platform to design and integrate a conversational user interface
- When an end-user writes or says something, referred to as an end-user expression, Dialogflow matches the end-user expression to the best intent in your agent

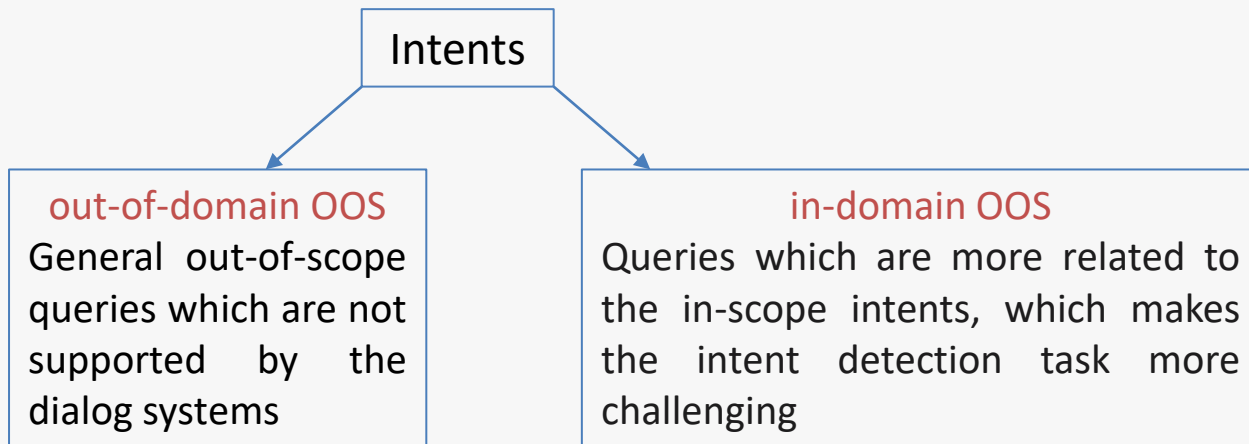
Used by : Dominos, Malaysian Airlines, KLM Royal Dutch Airlines, Verizon, CNN, etc.,





DATASET

- ❑ Motivation – NLU Benchmark (Few-Shot-Detection)
- ❑ Few-Shot-Intent-Detection is a repository designed for few-shot intent detection with/without Out-of-Scope (OOS) intents





DATASET

- The scope of our dataset has been inherited from 'CLINC150' intent dataset
- Our purpose was to built a chatbot for intent recognition we have developed a custom-built dataset

Features of Dataset:

```
    {"intents":  
      [{"tag": "greeting",  
        "patterns": ["Hi", "How are you", "Is anyone there?", "Hello", "Good day"],  
        "responses": ["Hello, thanks for visiting", "Good to see you again", "Hi there, how can I  
        help?"]  
      }  
    ]  
  }
```

Intents ←

User input ←

Bot responses ←



MODELS

☐ BERT

☐ DISTIL BERT

☐ LSTM



WHAT IS THE NEED FOR BERT?

- ❑ BERT uses bidirectional training
- ❑ It takes both the previous and next tokens into account simultaneously
- ❑ BERT applies the bidirectional training of Transformer to language modeling, learns the text representations



BERT

- ❑ Packages : Bert-for-tf2, BertModelLayer
- ❑ Pretrained model : uncased-bert-pretrained-model
- ❑ Tokenizer : FullTokenizer
- ❑ The pretrained bert model is embedded to the keras model with two dense layers and two dropout layers with value of 0.5 using the activation function of tanh and softmax

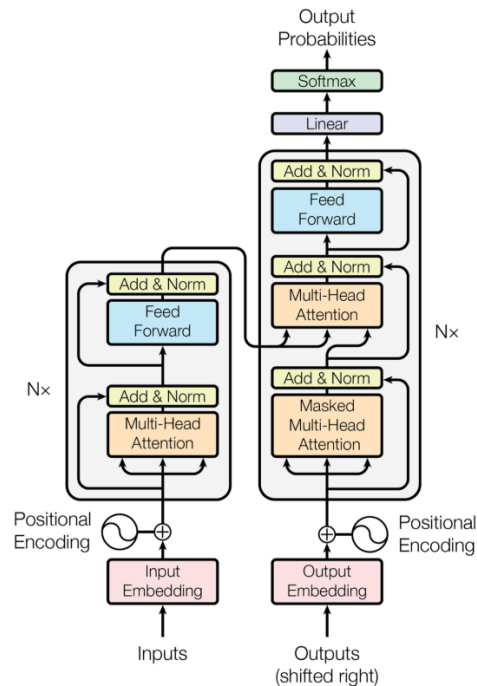


Figure 1: The Transformer - model architecture.



BERT

```
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_ids (InputLayer)	[(None, 16)]	0
bert (BertModelLayer)	(None, 16, 768)	108890112
lambda (Lambda)	(None, 768)	0
dropout (Dropout)	(None, 768)	0
dense (Dense)	(None, 768)	590592
dropout_1 (Dropout)	(None, 768)	0
dense_1 (Dense)	(None, 117)	89973
Total params: 109,570,677		
Trainable params: 109,570,677		
Non-trainable params: 0		

```
print("train acc", train_acc)  
print("test acc", test_acc)
```

```
train acc [4.759221453693032, 0.07734807]  
test acc [4.759221453693032, 0.07734807]
```

Accuracy is very poor

Reason: The training and test data of minimum 8000 entries respectively is required which was major drawback of our use case since it didn't had test set independently and was populated using train set itself



DISTILBERT

- ❑ Proposed to pre-train a smaller general-purpose language representation model
- ❑ The size of a BERT model by 40%, while retaining 97% of its language understanding capabilities and being 60% faster
- ❑ Distill BERT is smaller, faster and lighter model is cheaper to pre-train and demonstrates its capabilities for on-device computations in a proof-of-concept experiment and a comparative on-device study



DISTILBERT

	BERT	RoBERTa	DistilBERT	XLNet
Size (millions)	Base: 110 Large: 340	Base: 110 Large: 340	Base: 66	Base: ~110 Large: ~340
Training Time	Base: 8 x V100 x 12 days* Large: 64 TPU Chips x 4 days (or 280 x V100 x 1 days*)	Large: 1024 x V100 x 1 day; 4-5 times more than BERT.	Base: 8 x V100 x 3.5 days; 4 times less than BERT.	Large: 512 TPU Chips x 2.5 days; 5 times more than BERT.
Performance	Outperforms state-of-the-art in Oct 2018	2-20% improvement over BERT	3% degradation from BERT	2-15% improvement over BERT
Data	16 GB BERT data (Books Corpus + Wikipedia). 3.3 Billion words.	160 GB (16 GB BERT data + 144 GB additional)	16 GB BERT data. 3.3 Billion words.	Base: 16 GB BERT data Large: 113 GB (16 GB BERT data + 97 GB additional). 33 Billion words.
Method	BERT (Bidirectional Transformer with MLM and NSP)	BERT without NSP**	BERT Distillation	Bidirectional Transformer with Permutation based modeling



DISTILBERT

- ❑ For Distil BERT implementation we are making use of distilled-bert-uncased model
- ❑ The data is tokenized by using the pattern column of the data, and building a tensor out of the padded input and sending it to the distil bert model
- ❑ The model's performance is measured by using the classification report from sklearn.metrics package.

accuracy			0.32	37
macro avg	0.11	0.14	0.12	37
weighted avg	0.26	0.32	0.29	37



LSTM

- ❑ Introduced to avoid the long-term dependency problem
- ❑ LSTMs efficiently improve performance by memorizing the relevant information that is important and finding the pattern
- ❑ In LSTM we can use a multiple word string to find out the class to which it belongs
- ❑ Use of appropriate layers of embedding and encoding in LSTM, will be able to find out the actual meaning in input string and will give the most accurate output class



GLOVE

- ❑ It is an unsupervised learning algorithm developed by researchers at Stanford University aiming to generate word embeddings by aggregating global word co-occurrence matrices from a given corpus
- ❑ The basic idea behind the Glove word embedding is to derive the relationship between the words from statistics



LSTM

- ☐ Pretrained LSTM model is used
- ☐ GloVe : 6B.50d and 100d .txt
- ☐ Performed LSTM model building by embedding 1 dense layer with SoftMax activation function



LSTM

Model: "model_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 18)]	8
embedding_1 (Embedding)	(None, 18, 50)	200000050
lstm_1 (LSTM)	(None, 18, 128)	91648
dropout_1 (Dropout)	(None, 18, 128)	8
lstm_2 (LSTM)	(None, 128)	131584
dropout_2 (Dropout)	(None, 128)	8
dense_1 (Dense)	(None, 42)	5418
activation_1 (Activation)	(None, 42)	8

=====
Total params: 20,228,700
Trainable params: 228,658
Non-trainable params: 20,000,050
=====
Model: "model"

11/11 [=====] - 1s 6ms/step - loss: 0.2037 - accuracy: 0.9014
_loss: 0.2037387639284134 Accuracy: 0.9014492630958657

GloVe: 6B.50d
Accuracy : 0.90
Loss: 0.203

GloVe: 100d
Accuracy : 0.89
Loss: 0.2



WINNER

LSTM	0.904
BERT	0.077
DISTILBERT	0.32



GUI

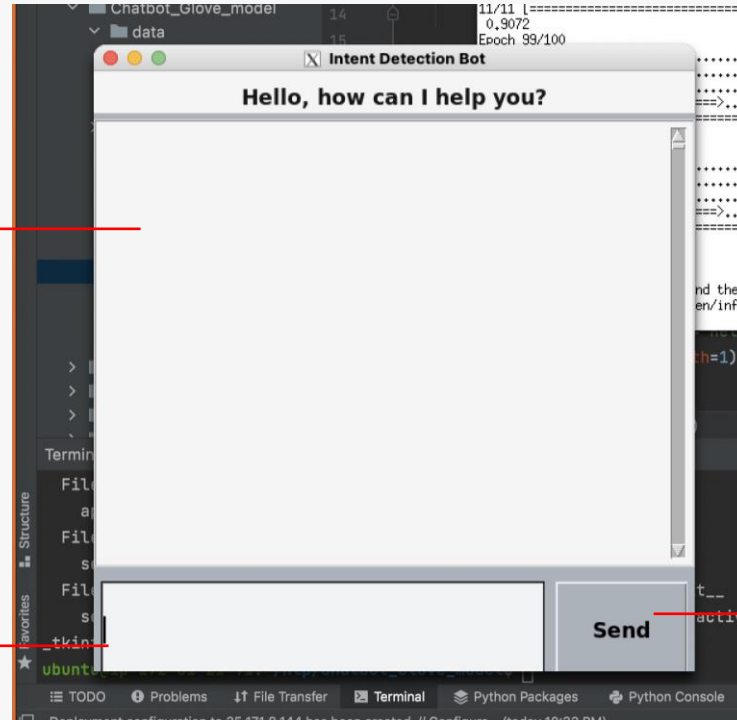
- ❑ Package : Tkinter
- ❑ To run the GUI, Xquartz has been used
- ❑ The GUI interacts with the user and retrieves the response using the best model **'LSTM'**



RESULTS

Conversation logs

Type your query



Press to send a message to the bot



RESULTS

A screenshot of a web application titled "Intent Detection Bot". The interface has a title bar with standard window controls. The main content area displays a conversation log. At the top, it says "Hello, how can I help you?". Below that, the user's input "You: Hi" is shown, followed by the bot's response "bot: Hi there, how can I help?". At the bottom of the chat area, there is a text input field and a "Send" button. The background shows a code editor with a file explorer on the left and a terminal at the bottom.

```
Chatbot_Glove_model 14 0,9072 Epoch 99/100
```

data

Intent Detection Bot

Hello, how can I help you?

You: Hi

bot: Hi there, how can I help?

Send

TODO Problems File Transfer Terminal Python Packages Python Console

Deployment configuration to 35.171.8.144 has been created. // Configure... (today 10:33 PM)

A screenshot of the same "Intent Detection Bot" web application, showing a continuation of the conversation. The user has entered "You: How to cancel the admission?" and the bot has responded with "bot: Cancellation of admission will be as per DTE rules.". The user then enters "You: bye" and the bot responds with "bot: See you!". The interface elements, including the title bar, chat area, input field, and bottom navigation bar, are consistent with the previous screenshot.

```
Chatbot_Glove_model 14 0,9072 Epoch 99/100
```

data

Intent Detection Bot

Hello, how can I help you?

You: Hi

bot: Hi there, how can I help?

You: How to cancel the admission?

bot: Cancellation of admission will be as per DTE rules.

You: bye

bot: See you!

Send

TODO Problems File Transfer Terminal Python Packages Python Console

Deployment configuration to 35.171.8.144 has been created. // Configure... (today 10:33 PM)



CONCLUSION & FUTURE SCOPE

- ❑ After making a comparison of all the model implementation we got to observe that LSTM is having the highest accuracy of 90%
- ❑ The further enhancements which can be done to this data is to increase the size of the train set and populate a test data
- ❑ Implementation of ConveRT - A dual sentence encoder , it is effective, affordable, and quick to train also the size of the ConveRT model is less compared to the BERT model



REFERENCES

- <https://arxiv.org/abs/1805.10190>
- <https://github.com/sonos/nlu-benchmark/tree/master/2017-06-custom-intent-engines>
- <https://github.com/huggingface/transformers>
- <https://paperswithcode.com/task/intent-detection>
- <https://www.sciencedirect.com/science/article/pii/S1877050918320374>
- <https://analyticsindiamag.com/hands-on-guide-to-word-embeddings-using-glove/>



PROJECT REPO

❑ The project implementation and codes can be found on the following repo:

<https://github.com/Rehamanikandan/Final-Project-Group6>